

```
In [ ]: import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
import statsmodels.api as sm
import matplotlib.pyplot as plt
import seaborn as sb
import plotly.express as px
plt.style.use('default')
```

### 1. Read dataset

```
In [ ]: boston = pd.read_csv("D:\\PROGRAMMING\\Datasets\\Boston.csv")
boston.head()
```

```
Out[ ]:
```

	Unnamed: 0	crim	zn	indus	chas	nox	rm	age	dis	rad	tax	ptratio	black	lstat	medv
0	1	0.00632	18.0	2.31	0	0.538	6.575	65.2	4.0900	1	296	15.3	396.90	4.98	24.0
1	2	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	242	17.8	396.90	9.14	21.6
2	3	0.02729	0.0	7.07	0	0.469	7.185	61.1	4.9671	2	242	17.8	392.83	4.03	34.7
3	4	0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	3	222	18.7	394.63	2.94	33.4
4	5	0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	3	222	18.7	396.90	5.33	36.2

### 2. Find dependent and independent variables

```
In [ ]: X = pd.DataFrame(boston.iloc[:, :-1])
y = pd.DataFrame(boston.iloc[:, -1])
```

### 3. Check for the significance

```
In [ ]: # The inclusion of a constant allows the regression line to have an intercept point with
# the y-axis, even when all independent variables are zero.
```

```
In [ ]: ##level of significance
alpha = 0.05

## Add constant to the independent variable
X = sm.add_constant(X)

sig_est = sm.OLS(y, X)## OLS( Ordinary Least Square )
result = sig_est.fit()
print(result.summary())
```

```

=====
                        OLS Regression Results
=====
Dep. Variable:          medv    R-squared:                0.741
Model:                  OLS      Adj. R-squared:            0.734
Method:                 Least Squares    F-statistic:          100.6
Date:                  Tue, 11 Jul 2023    Prob (F-statistic):    3.44e-134
Time:                  19:26:22    Log-Likelihood:        -1498.0
No. Observations:      506    AIC:                   3026.
Df Residuals:          491    BIC:                   3089.
Df Model:              14
Covariance Type:       nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
const	36.4614	5.101	7.148	0.000	26.439	46.484
Unnamed: 0	-0.0025	0.002	-1.215	0.225	-0.007	0.002
crim	-0.1088	0.033	-3.310	0.001	-0.173	-0.044
zn	0.0480	0.014	3.484	0.001	0.021	0.075
indus	0.0199	0.061	0.324	0.746	-0.101	0.141
chas	2.7052	0.861	3.141	0.002	1.013	4.398
nox	-17.5416	3.822	-4.589	0.000	-25.052	-10.031
rm	3.8392	0.418	9.175	0.000	3.017	4.661
age	-0.0019	0.013	-0.145	0.885	-0.028	0.024
dis	-1.4933	0.200	-7.471	0.000	-1.886	-1.101
rad	0.3249	0.068	4.771	0.000	0.191	0.459
tax	-0.0116	0.004	-3.046	0.002	-0.019	-0.004
ptratio	-0.9480	0.131	-7.246	0.000	-1.205	-0.691
black	0.0094	0.003	3.485	0.001	0.004	0.015
lstat	-0.5262	0.051	-10.377	0.000	-0.626	-0.427

```

=====
Omnibus:                175.545    Durbin-Watson:          1.084
Prob(Omnibus):           0.000    Jarque-Bera (JB):        760.925
Skew:                    1.502    Prob(JB):                5.85e-166
Kurtosis:                8.202    Cond. No.                1.68e+04
=====

```

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 1.68e+04. This might indicate that there are strong multicollinearity or other numerical problems.

```

In [ ]: ##Checking for the simple linear regression
check = sm.OLS(boston['medv'] , boston['age']).fit()
check.summary()

```

```

Out[ ]:
                        OLS Regression Results
=====
Dep. Variable:          medv    R-squared (uncentered):    0.644
Model:                  OLS      Adj. R-squared (uncentered):    0.644
Method:                 Least Squares    F-statistic:          915.1
Date:                  Tue, 11 Jul 2023    Prob (F-statistic):    1.85e-115
Time:                  19:26:22    Log-Likelihood:        -2071.5
No. Observations:      506    AIC:                   4145.
Df Residuals:          505    BIC:                   4149.
Df Model:              1
Covariance Type:       nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
age	0.2636	0.009	30.250	0.000	0.246	0.281

```

=====
Omnibus: 27.739    Durbin-Watson: 0.357
Prob(Omnibus): 0.000    Jarque-Bera (JB): 19.564
Skew: 0.369    Prob(JB): 5.65e-05
Kurtosis: 2.380    Cond. No. 1.00
=====

```

Notes:

- [1]  $R^2$  is computed without centering (uncentered) since the model does not contain a constant.
- [2] Standard Errors assume that the covariance matrix of the errors is correctly specified.

#### 4. Dropping the necessary columns

```
In [ ]: ## unnamed:0, Indus and Age are statistically insignificant for we accept null hypothesis in this
## case and reject the alternate hypothesis, and we will drop those attributes from our model

X = X.drop(["Unnamed: 0", "indus", "age"], axis = 1)
X.head()
```

```
Out[ ]:
```

	const	crim	zn	chas	nox	rm	dis	rad	tax	ptratio	black	lstat
0	1.0	0.00632	18.0	0	0.538	6.575	4.0900	1	296	15.3	396.90	4.98
1	1.0	0.02731	0.0	0	0.469	6.421	4.9671	2	242	17.8	396.90	9.14
2	1.0	0.02729	0.0	0	0.469	7.185	4.9671	2	242	17.8	392.83	4.03
3	1.0	0.03237	0.0	0	0.458	6.998	6.0622	3	222	18.7	394.63	2.94
4	1.0	0.06905	0.0	0	0.458	7.147	6.0622	3	222	18.7	396.90	5.33

5. Find the coefficients | p\_values | Confidence Interval

```
In [ ]: print("\nThe Coefficiencts are : ")
result.params
```

The Coefficiencts are :

```
Out[ ]: const          36.461352
Unnamed: 0          -0.002526
crim                -0.108762
zn                  0.048031
indus               0.019932
chas                2.705245
nox                -17.541602
rm                  3.839225
age                -0.001938
dis                -1.493304
rad                 0.324925
tax                -0.011598
ptratio            -0.947985
black               0.009357
lstat              -0.526184
dtype: float64
```

```
In [ ]: print("The P-Values are: ")
result.pvalues
```

The P-Values are:

```
Out[ ]: const          3.209691e-12
Unnamed: 0          2.250457e-01
crim                1.000250e-03
zn                  5.375059e-04
indus               7.458713e-01
chas                1.785946e-03
nox                 5.658365e-06
rm                  1.245587e-18
age                 8.848664e-01
dis                 3.682773e-13
rad                 2.426287e-06
tax                 2.443267e-03
ptratio            1.670700e-12
black               5.364596e-04
lstat              6.050328e-23
dtype: float64
```

```
In [ ]: print("Confidence Intervals are: ")
result.conf_int()
```

Confidence Intervals are:

Out[ ]:	0	1
<b>const</b>	26.438882	46.483822
<b>Unnamed: 0</b>	-0.006612	0.001560
<b>crim</b>	-0.173316	-0.044209
<b>zn</b>	0.020946	0.075115
<b>indus</b>	-0.100841	0.140705
<b>chas</b>	1.012960	4.397531
<b>nox</b>	-25.051861	-10.031344
<b>rm</b>	3.017106	4.661344
<b>age</b>	-0.028227	0.024350
<b>dis</b>	-1.886054	-1.100554
<b>rad</b>	0.191101	0.458750
<b>tax</b>	-0.019078	-0.004117
<b>ptratio</b>	-1.205026	-0.690945
<b>black</b>	0.004081	0.014632
<b>lstat</b>	-0.625808	-0.426560

## 6. Train Test Split

```
In [ ]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random_state = 100)
print(X_train.shape)
print(X_test.shape)
print(y_train.shape)
print(y_test.shape)
```

(404, 12)  
 (102, 12)  
 (404, 1)  
 (102, 1)

## 7. Training Our Model

```
In [ ]: mlr = LinearRegression()
mlr.fit(X_train, y_train)
```

```
Out[ ]: ▼ LinearRegression
LinearRegression()
```

## 8. Predicting Value

```
In [ ]: y_predict = mlr.predict(X_test)

final = pd.DataFrame({'Actual':y_test.values.flatten(), 'Predicted':y_predict.flatten()})
final
```

Out[ ]:

	Actual	Predicted
0	34.6	34.496490
1	31.5	30.868682
2	20.6	22.304769
3	14.5	18.131193
4	16.2	20.541658
...	...	...
97	50.0	36.370316
98	7.2	18.015547
99	50.0	23.490485
100	14.0	13.702219
101	11.0	14.314579

102 rows × 2 columns

```
In [ ]: px.scatter(final, 'Actual', 'Predicted', trendline = 'ols', trendline_color_override='blue')
```

9. Necessary observations

```
In [ ]: mae = mean_absolute_error(y_test, y_predict)
mse = mean_squared_error(y_test, y_predict)
rmse = np.sqrt(mse)
r2 = r2_score(y_test, y_predict)

n = len(y_test)      ## no of samples
p = X_test.shape[1]  ## no of predictors
adjusted_r2 = 1 - ( 1 - r2 )*( n - 1) / n - p - 1

print("mean_absolute_error is: ", mae)
print("mean_squared_error is : ", mse)
print("root_mean_squared_error is : ", rmse)
print("r square is : ", r2)
print("adjusted r square is : ", adjusted_r2)## it decreases as features/predictors increase
```

mean\_absolute\_error is: 3.2518545636225586  
mean\_squared\_error is : 23.425938278313655  
root\_mean\_squared\_error is : 4.840034945980623  
r square is : 0.7574812283240356  
adjusted r square is : -12.240141136659533